

MODEER

Product & engineering handover

A personal AI assistant and nine specialists that share what they know about one person.

WORKING TREE · UNCOMMITTED

Evidence as of 2026-09-15

HEAD 1ed17eb · branch fixes/launch-readiness-20260911



Modeer — character art from frontend/public/art/sunshine

Not approved for public production.

Read every claim with its evidence label

IMPLEMENTED

Present in source code; not proof of production behaviour.

VERIFIED

A dated check passed in the named environment only.

HISTORICAL

Done earlier under different code, prompts or rubric.

PROPOSED

Planned work. Nothing is built.

UNRESOLVED DEFECT

Measured problem that is still open.

COVERAGE GAP

Behaviour nobody has tested yet.

OPERATIONAL GATE

Needs the real host, storage or a person to close.

OWNER DECISION

Only the owner can decide.

Baseline: HEAD 1ed17eb plus an uncommitted working tree. The Sept 13–14 fixes are implemented and locally verified, not released.

One persistent personal context, ten conversations

For	Does	Does not
• One person who wants help across study, work, money, health habits and writing. • Private, self-hosted deployment run by an operator.	• Chats with Modeer or a specialist; replies stream live. • Remembers facts shared with the team, plus private notes per specialist. • Tracks goals and a daily Today list.	• Browse, send email, book, buy or read calendars. • Offer mobile apps, payments or push reminders.

Release status · 2026-09-15			Differentiator: one personal-context layer and ten human characters, not ten unrelated chatbots.
Local use	Usable with known limits	VERIFIED · LOCAL	
Invite-only beta	On hold until gates close	OPERATIONAL GATE	
Public production	On hold	OPERATIONAL GATE	

Modeer and nine specialists

Artwork: AI-generated 2026-09-10 · no licence inventory yet



Modeer

Assistant & context keeper



Study

Learning coach




Career

CVs & interviews




Research

Analysis



Writing

Drafting & editing




Travel

Itineraries



Shopping

Choices & trade-offs



Finance

Budgeting



Fitness


Training plans




Email

Email drafts

The user journey, end to end




When a reply does not finish




Limits: no email reset or email verification; Today does not read a calendar; Continue and Try again are user clicks, never automatic; a reload fetches the server's record as the source of truth.

What it looks like today



Home — hero composer and featured specialists


LOCAL DEV · MOCK MODEL · SYNTHETIC DATA



Memory — shared facts (synthetic) and private notes by specialist

- Also captured: Team, Goals, Account, Login, chat, 390px views (appendix).
- Chat text on these screens is the mock provider's placeholder, not model output.

Sunshine & Ink, measured





Ink on pink · 5.5:1

Was white on pink at 3.27:1. Fixed in 80cf49a.


Accessibility evidence

- 25 automated scans, 0 findings — desktop + 390px, local HTTPS rehearsal
- Keyboard pass — 11 routes, visible focus, signup to first message
- **Not done:** physical phones, human screen-reader pass




390px in desktop Chrome — not a physical phone

A modular monolith behind one HTTPS proxy




One chat turn, from click to memory



SSE events: start · delta · end · memory · error

Watch out: "end" is not "memory saved". No replay, no idempotency key — reload history after an uncertain stream.

Every reply records how it ended



Refunds: a provider failure before any text is refunded; partial or cancelled replies stay charged.

Agents are configuration, not services

Agent	Model	Temp	Output cap	Prompt v
Modeer	gpt-oss-120b	0.55	650	4
Study	gpt-oss-120b	0.5	800	5
Career	gpt-oss-120b	0.55	650	5
Research	gpt-oss-120b	0.4	800	6
Writing	qwen3.8-27b	0.3	1600	5
Travel	gpt-oss-120b	0.6	800	8
Shopping	gpt-oss-120b	0.5	1100	5
Finance	gpt-oss-120b	0.4	800	4
Fitness	gpt-oss-120b	0.5	800	6
Email	gpt-oss-120b	0.6	1000	3

How an agent is built


- config.py — model, caps, framework, safety
- prompt.md — role instructions
- evals.json — quality cases
- No automatic router: the person picks who to talk to.

Ask My Team

IMPLEMENTED · OFF

- Off by default (TEAM_ENABLED=false); no screen calls it.
- If on: up to 5 specialists, shared context only, no learning, then Modeer synthesises — up to 6 model calls.

Who can see which fact



Tested at the provider's input with positive controls. Sensitive-fact filtering is best effort; the Account switch stops new learning, not existing context.

Ten tables, one owner per row

Account

users
recovery_codes
usage_buckets

Conversation

conversations
messages
agents

Personal context

shared_memories
agent_memories
goals
briefings

Verified on disposable PostgreSQL 16 · 2026-09-14

- Empty → 0004 and 0004 → 0003 → 0004 passed; a sentinel user survived.
- **Downgrade deletes password and recovery data.** Roll back with a tested backup.

Learn

extract · validate

Store

user edits protected

Export

JSON, no secrets

Delete

cascade · clears session

Backups

not erased by deletion

Migrations

0001 base tables
0002 usage buckets
0003 session epoch
0004 passwords,
memory switch, recovery
codes

Revocable sessions, atomic credential changes

Session	Credentials	Limits
• Signed cookie, HttpOnly, SameSite=Strict, Secure • Epoch bump ends every session for one account • 7-day default lifetime	• scrypt passwords, 10–128 chars • 10 single-use recovery codes, hashed • Atomic recovery (working tree)	• 6 chats / account / min • 60,000 budget units / UTC day • 10 sign-in tries / email / 15 min

VERIFIED · POSTGRESQL 16 · LOCAL

Recovery races: the same code twice gives one success and one rejection; two different codes give sessions at epochs 1 and 2.

Known limits: no email verification or reset; no per-device sessions; throttles untested against distributed abuse; every account shares one free-tier provider quota.

Controls exist; the real host still has to prove them

Security controls

- Nonce CSP + strict headers
- Safe Markdown links
- API docs off in production
- Logs without message content
- Locked, audited dependencies

Backups & alerts

- Encrypted nightly backup script
- Restore check after every backup
- Watchdog alerts on state change
- Webhook or Telegram channel

Operator must

OPERATIONAL GATE

- Install the cron jobs
- Provide truly remote storage
- Confirm a person receives alerts
- Audit legacy memories

VERIFIED · DISPOSABLE DOCKER + POSTGRESQL 16

10 operations checks: schema-0004 restore matched every table, including credentials and consent.

Not proven: the copy sat in a second local folder; alerts reached a synthetic receiver, not a person.

How it got here

Labels: historical → committed → uncommitted → locally verified



Commit or otherwise preserve the working tree before anything else: the latest security and data fixes exist only there.

Five kinds of evidence — never add them up

351 backend tests	23 HTTPS browser checks	25 accessibility scans · 0 findings	10 operations checks	13 live journey checks
WHERE Locked production image · 2 races on PostgreSQL 16, rest SQLite	WHERE Production images on localhost · scripted model	WHERE Desktop Chrome, desktop + 390px widths	WHERE Disposable Docker + PostgreSQL 16	WHERE Real model via in-process test client
LIMIT Not load or multi-worker	LIMIT Not the public domain or real model	LIMIT Not phones or a screen reader	LIMIT Copy was local; alerts synthetic	LIMIT Not a real network disconnect

Answer quality is the main open defect

UNRESOLVED DEFECT

Personalization (3 samples)	Pass	Fail	Ungraded	What went wrong (manual review)
Modeer · Study · Research · Email	3 each	0	0	—
Shopping	0	3	0	Unverified budget fit; invented \$; “latest” claim
Travel	1	2	0	Asserted departure times; residence ≠ nationality
Career	1	2	0	Invented visa and mentorship; unsupported timeline
Writing	1	2	0	Very short bios; missed a supplied aspiration
Finance	2	1	0	Unsourced current property-price range
Fitness	1	1	1	Equipment not provided; judge JSON unparsable
Unknown-fact cases (Study, Career, Writing)	9	0	0	—

Automated totals 27 · 11 · 1 are not a quality score. Manual checks found judge false positives and missed defects. 33/39 on Sep 13 is historical (older prompts and rubric).

Three release levels, three different answers

Local use	Invite-only beta	Public production
Usable	Hold	Hold
Requires / known • Synthetic or non-critical data • Known quality defects • Fixes still uncommitted	Before it can open • Real domain + HTTPS staging • Remote backup, scheduler, restore • Alerts reach a person • Quality decision from measured failures	Before it can open • All beta gates • Signup abuse and quota capacity • Support & privacy ownership • Physical phones + screen reader • Asset rights handover

What is still open, by kind

UNRESOLVED DEFECT

Unresolved defects

- Grounding errors in 6 of 10 personalization areas
- Judge JSON unreliable; judge misgrades both ways

COVERAGE GAP

Coverage gaps

- Screen reader and physical phones
- Load, multi-worker, real disconnects
- No CI pipeline; no stream replay

OPERATIONAL GATE

Operational gates

- Real domain, TLS, staging
- Remote backup + scheduler + restore
- Alert receipt; legacy memory audit

OWNER DECISION

Owner decisions

- Signup mode, memory default
- Quality bar and provider capacity
- Retention, support, asset rights

Your first ten moves

1

Preserve and commit the working tree with the owner

2

Reproduce all local checks

3

Remove only the old rehearsal Docker project

4

Collect owner decisions

5

Fix and re-measure answer quality

6

Stage on the real HTTPS domain

7

Remote backups, scheduler, restore

8

Confirm a person gets alerts

9

Phones, screen reader, memory audit → beta call

10

Then: mobile feasibility spike

Appendix

A1 API surface

A2 Configuration (names only)

A3 Reproducible commands

A4 More screens

A5 Evidence index and document discrepancies

A1 · API surface (all under /api)

Area	Endpoints	Notes
Auth	status · login · signup · recover · logout · account · password · recovery-codes · sign-out-everywhere	Signup gated by SIGNUP_ENABLED
User	users/me (GET, PATCH) · users/me/export · users/me/delete	Delete needs confirm: DELETE
Agents & chat	agents · agents/{id} · agents/{id}/chat/stream (SSE)	Sync chat is 404 in production
Conversations	conversations (GET, POST) · conversations/{id} (GET, DELETE)	Ownership checked
Memory	memory/shared[/]{id}] · memory/agent/{agent_id} · memory/agent[/]{id}]	Explicit null → 422
Goals & Today	goals[/]{id}] · briefings/today	No calendar integration
Team	team/ask	404 while TEAM_ENABLED=false
Health & admin	health · health/detail · legal/privacy · admin/dashboard · admin/metrics · admin/test-alert	Admin = ADMIN_ACCOUNTS

A2 · Configuration — names only, never values

App & auth	APP_NAME · ENVIRONMENT · DEBUG · AUTH_REQUIRED · AUTH_SECRET · AUTH_ACCESS_KEYS · SESSION_SECONDS · SIGNUP_ENABLED · TEAM_ENABLED · ADMIN_ACCOUNTS
Web & data	DOMAIN · DB_PASSWORD · DATABASE_URL · FRONTEND_URL · NEXT_PUBLIC_API_URL · BACKEND_URL
Model & memory	LLM_PROVIDER · LLM_API_KEY · LLM_MODEL · LLM_BASE_URL · LLM_REASONING_EFFORT · LLM_TIMEOUT_SECONDS · LLM_MAX_TOKENS · LLM_TEMPERATURE · MEMORY_EXTRACTION · MEMORY_TIMEOUT_SECONDS · MEMORY_STORE_SENSITIVE · MEMORY_MIN_CONFIDENCE
Limits & alerts	ACCOUNT_REQUESTS_PER_MINUTE · ACCOUNT_DAILY_TOKEN_BUDGET · ALERT_WEBHOOK_URL · ALERT_TELEGRAM_BOT_TOKEN · ALERT_TELEGRAM_CHAT_ID · ALERT_ERROR_THRESHOLD
Backups & watchdog	MODEER_COMPOSE_FILE · MODEER_BACKUP_DIR · MODEER_BACKUP_OFFHOST · MODEER_BACKUP_KEEP_DAYS · MODEER_BACKUP_PASSPHRASE_FILE · MODEER_URL · MODEER_WATCHDOG_STATE · MODEER_WATCHDOG_TIMEOUT

A3 · Reproduce the checks

```
# backend
.venv/Scripts/python.exe -m pytest -q
.venv/Scripts/python.exe -m ruff check .
.venv/Scripts/python.exe -m pip check

# frontend
npm ci
npm run typecheck
npm run lint
npm run build
node tools/links-check.cjs
node tools/stream-check.cjs
node tools/api-check.cjs
npm audit --omit=dev
```


Local production rehearsal

- deploy/tests/production-rehearsal/ README
- Fresh Compose project name each run
- Seed users from PowerShell


Live quality (spends free quota)

- quality_check.py --samples 3
- tools.rejudge_quality
- journey_check.py --db <new>
- Freeze sources; new output names


A4 · More screens from the working tree




Team




Goals — synthetic goals



Account



Chat — mock provider placeholder



Login

A5 · Evidence index and discrepancies

Artifact (2026-09-14 unless noted)	Shows	Limit
production-verification.json	Image digests · 351 tests · migrations	Local images
production-rehearsal.json	23 HTTPS browser checks	Scripted model
accessibility-rehearsal.json	25 scans · keyboard stops	No phones / screen reader
operations-rehearsal.json	10 backup and alert checks	Local copy · synthetic alerts
quality-remediation-14b.json	39 answers · 27/11/1	Judge imperfect
journey-rehearsal.txt	13 live journey checks	TestClient closure

Document discrepancies

- PROJECT-REPORT \$5 prompt versions are older than the code
- DEPLOYMENT.md: “20 checks” and contrast note are stale
- HANDOFF.md, STATUS.md stop at Sep 13
- Older notes said 11 tables; there are 10
- Rehearsal Docker cleanup not verified